

Games

Multi-agent interaction

Up to this point:

This week:

Multi-agent interaction

Up to this point:

- Stochastic Uncertainty

This week:

Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

- Interaction Uncertainty

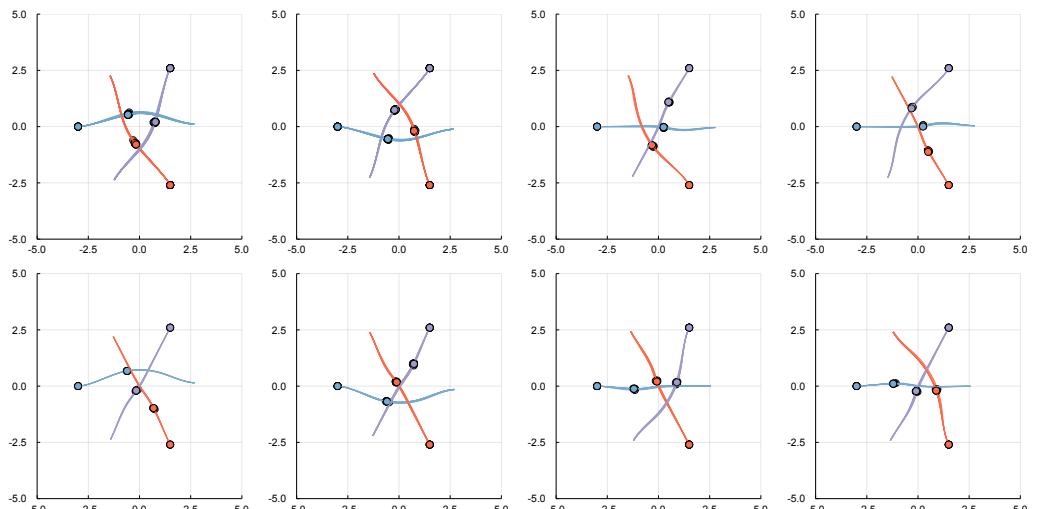
Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

- Interaction Uncertainty
- Game Solution Concepts



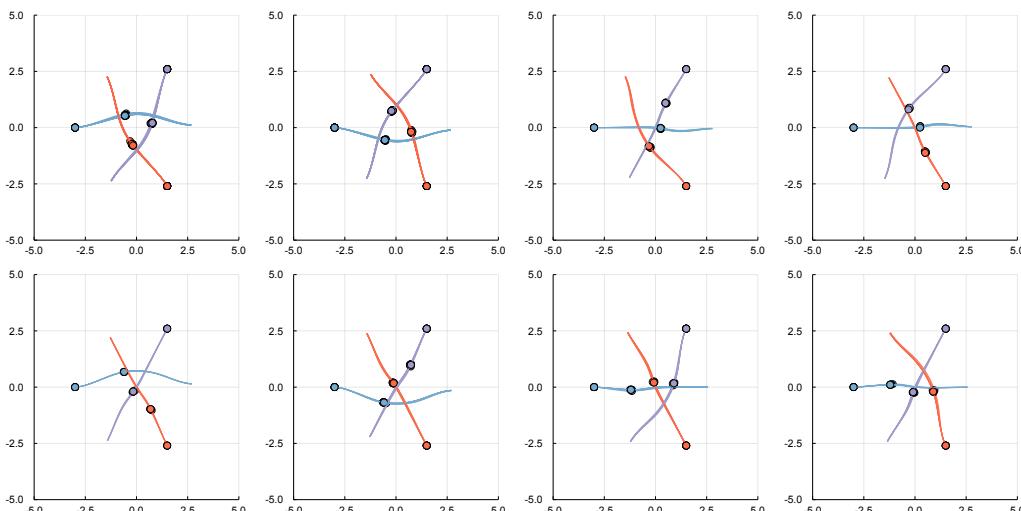
Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

- Interaction Uncertainty
- Game Solution Concepts



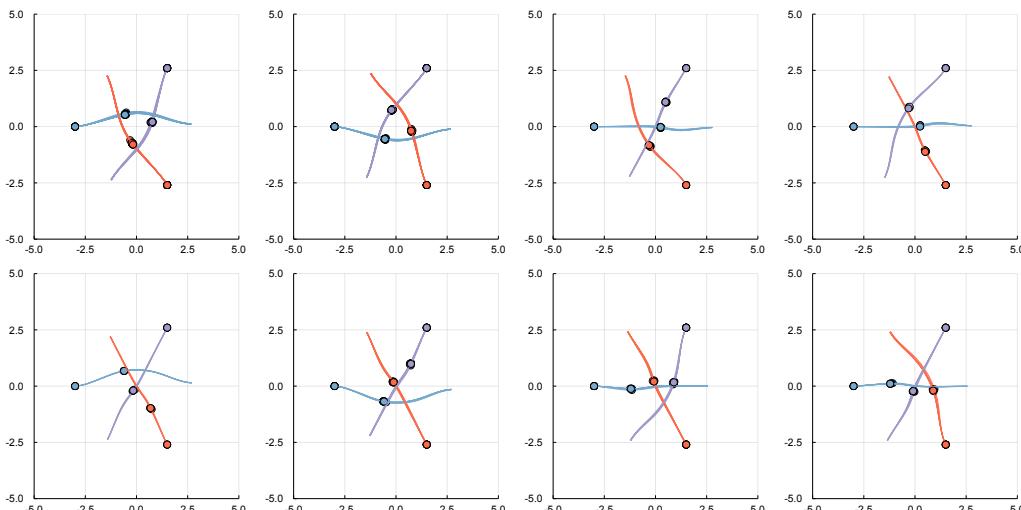
Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

- Interaction Uncertainty
- Game Solution Concepts



Kriegspiel

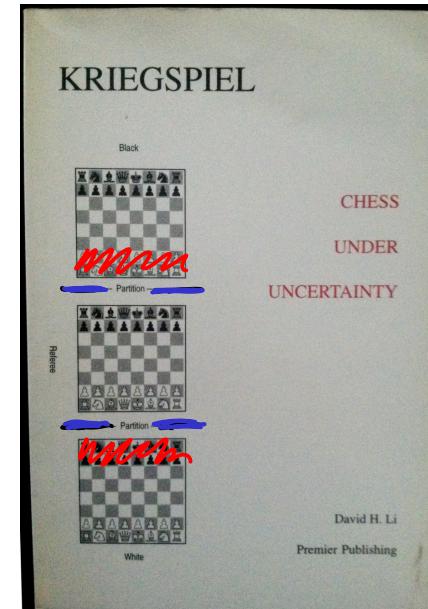
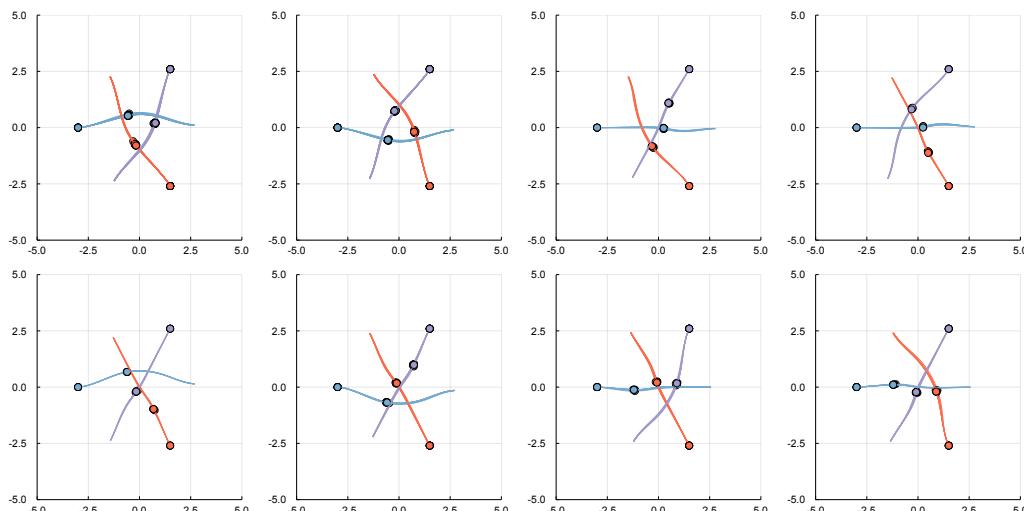
Multi-agent interaction

Up to this point:

- Stochastic Uncertainty
- Optimization

This week:

- Interaction Uncertainty
- Game Solution Concepts



Kriegspiel

Outline

Outline

1. Matrix Games

1. Linear Programming
2. Regret Matching

Outline

1. Matrix Games
 1. Linear Programming
 2. Regret Matching
2. Turn-taking Games
 1. Minimax Trees
 2. α - β Pruning
 3. Monte Carlo Tree Search

Outline

1. Matrix Games
 1. Linear Programming
 2. Regret Matching
2. Turn-taking Games
 1. Minimax Trees
 2. α - β Pruning
 3. Monte Carlo Tree Search
3. Imperfect Information Games
 1. Extensive-form Games
 2. Counterfactual Regret Minimization

Outline

1. Matrix Games
 1. Linear Programming
 2. Regret Matching
2. Turn-taking Games
 1. Minimax Trees
 2. α - β Pruning
 3. Monte Carlo Tree Search
3. Imperfect Information Games
 1. Extensive-form Games
 2. Counterfactual Regret Minimization

Sources:

- *AI: A Modern Approach* by Russel and Norvig
- "An Introduction to Counterfactual Regret Minimization in Games" by Neller and Lanctot

Game Solution Concepts: Nash and Correlated Equilibria

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Battle of the Sexes

		M	G	
		M	$2, 1$	$0, 0$
Monica	M	$2, 1$	$0, 0$	
	G	$0, 0$	$1, 2$	

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

Definitions:

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Battle of the Sexes

Gary

	M	G
M	$2, 1$	$0, 0$
G	$0, 0$	$1, 2$

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Definitions:

- Strategy/Policy: Probability distribution of actions for a player

Battle of the Sexes

		Gary	
		M	G
Monica	M	$2, 1$	$0, 0$
	G	$0, 0$	$1, 2$

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Battle of the Sexes

		Gary	
		M	G
Monica	M	$2, 1$	$0, 0$
	G	$0, 0$	$1, 2$

Definitions:

- Strategy/Policy: Probability distribution of actions for a player
- Strategy Profile: Set of strategies for all players

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	<i>R</i>	<i>P</i>	<i>S</i>
<i>R</i>	0, 0	-1, 1	1, -1
<i>P</i>	1, -1	0, 0	-1, 1
<i>S</i>	-1, 1	1, -1	0, 0

Battle of the Sexes

Gary

	<i>M</i>	<i>G</i>
<i>M</i>	2, 1	0, 0
<i>G</i>	0, 0	1, 2

Definitions:

- Strategy/Policy: Probability distribution of actions for a player
- Strategy Profile: Set of strategies for all players
- Nash Equilibrium: A strategy where there is no incentive for any player to unilaterally change strategy

Game Solution Concepts: Nash and Correlated Equilibria

Rock Paper Scissors

	<i>R</i>	<i>P</i>	<i>S</i>
<i>R</i>	0, 0	-1, 1	1, -1
<i>P</i>	1, -1	0, 0	-1, 1
<i>S</i>	-1, 1	1, -1	0, 0

Battle of the Sexes

		<i>M</i>	<i>G</i>	
		<i>M</i>	2, 1	0, 0
<i>M</i>	<i>M</i>	2, 1	0, 0	
	<i>G</i>	0, 0	1, 2	

Definitions:

- Strategy/Policy: Probability distribution of actions for a player
- Strategy Profile: Set of strategies for all players
- Nash Equilibrium: A strategy where there is no incentive for any player to unilaterally change strategy
- Correlated Equilibrium: Strategy profile with a shared random source, where no player has an incentive to deviate

Calculating Mixed-Strategy Equilibria

Calculating Mixed-Strategy Equilibria

Gary

	M	G
Monica	M	2, 1
	G	0, 0

	M	G
Monica	M	2, 1
	G	0, 0

	M	G
Monica	M	2, 1
	G	0, 0

	M	G
Monica	M	2, 1
	G	0, 0

Calculating Mixed-Strategy Equilibria

Gary

	M	G
Monica	M	2, 1
	G	0, 0

Monica	M	G
	2, 1	0, 0
	0, 0	1, 2

By hand

Calculating Mixed-Strategy Equilibria

Gary

	M	G	
Monica	M	$2, 1$	$0, 0$
	G	$0, 0$	$1, 2$

By hand

$$\sigma_{\text{Gary}}(M) = x$$

Calculating Mixed-Strategy Equilibria

Gary

	<i>M</i>	<i>G</i>
Monica	<i>M</i>	2, 1
	<i>G</i>	0, 0

Monica	<i>M</i>	0, 0
	<i>G</i>	1, 2

By hand

$$\sigma_{\text{Gary}}(M) = x$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, M) = 2x,$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, G) = 1 - x$$

Calculating Mixed-Strategy Equilibria

Gary

	M	G
Monica	M	2, 1
	G	0, 0

	M	G
Monica	M	2, 1
	G	0, 0

By hand

$$\sigma_{\text{Gary}}(M) = x \frac{1}{3}$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, M) = 2x, \frac{2}{3}$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, G) = 1 - x \frac{2}{3}$$

If $x = 1/3$, these are equal, so Monica is indifferent. By symmetry, $\sigma = \{(\frac{2}{3}, \frac{1}{3}), (\frac{1}{3}, \frac{2}{3})\}$ is a Nash Equilibrium.

Calculating Mixed-Strategy Equilibria

Gary

	M	G
Monica	M	2, 1
	G	0, 0

	M	G
Monica	M	2, 1
	G	0, 0

By hand

$$\sigma_{\text{Gary}}(M) = x$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, M) = 2x,$$

$$U_{\text{Monica}}(\sigma_{\text{Gary}}, G) = 1 - x$$

If $x = 1/3$, these are equal, so Monica is indifferent. By symmetry, $\sigma = \{(\frac{2}{3}, \frac{1}{3}), (\frac{1}{3}, \frac{2}{3})\}$ is a Nash Equilibrium.

$$\underset{\pi^U}{\text{minimize}} \quad \sum_i (U^i - U^i(\pi))$$

$$\text{subject to } U^i \geq U^i(a^i, \pi^{-i}) \text{ for all } i, a^i$$

$$\begin{cases} \sum_{a^i} \pi^i(a^i) = 1 \text{ for all } i \\ \pi^i(a^i) \geq 0 \text{ for all } i, a^i \end{cases}$$

Regret Matching

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(\underline{a'_i}, \underline{a_{-i}}) - \underline{U(\mathbf{a})}$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

For some number of iterations:

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

For some number of iterations:

1. Compute a regret matching profile ($\sigma \propto$ normalized cumulative regret)

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

For some number of iterations:

1. Compute a regret matching profile ($\sigma \propto$ normalized cumulative regret)
2. Add strategy profile to profile sum

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

For some number of iterations:

1. Compute a regret matching profile ($\sigma \propto$ normalized cumulative regret)
2. Add strategy profile to profile sum
3. Select an action

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

For some number of iterations:

1. Compute a regret matching profile ($\sigma \propto$ normalized cumulative regret)
2. Add strategy profile to profile sum
3. Select an action
4. Compute regrets

Regret Matching

For action profile \mathbf{a} , player i 's *regret* for not playing action a' is:

$$R(a'_i) = U(a'_i, a_{-i}) - U(\mathbf{a})$$

that is, the utility missed out on by not playing a' .

Positive regret means that action would have been better.

Regret Matching:

$$\bar{\sigma} \quad \bar{r}$$

For some number of iterations:

1. Compute a regret matching profile ($\sigma \propto$ normalized cumulative regret)
2. Add strategy profile to profile sum $\bar{\sigma}$
3. Select an action from $\bar{\sigma}$
4. Compute regrets \bar{r}
5. Add regrets to cumulative regrets \bar{F}

Turn-taking Games

Turn-taking Games



Terminology

Turn-taking Games



Terminology

- Max and Min players

Turn-taking Games



Terminology

- Max and Min players
- deterministic

Turn-taking Games



Terminology

- Max and Min players
- deterministic
- two player

Turn-taking Games



Terminology

- Max and Min players
- deterministic
- two player
- zero-sum

Turn-taking Games



Terminology

- Max and Min players
- deterministic
- two player
- zero-sum
- perfect information

Minimax Trees

Minimax Trees

MDP Expectimax Tree

Minimax Trees

MDP Expectimax Tree

Minimax Tree

Minimax Trees

MDP Expectimax Tree

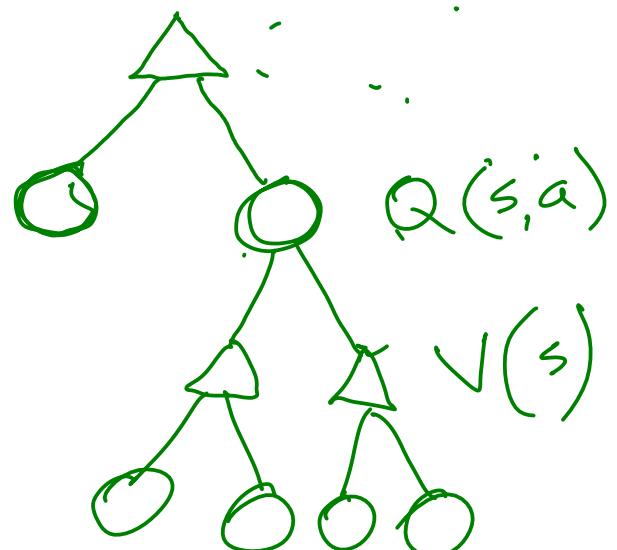
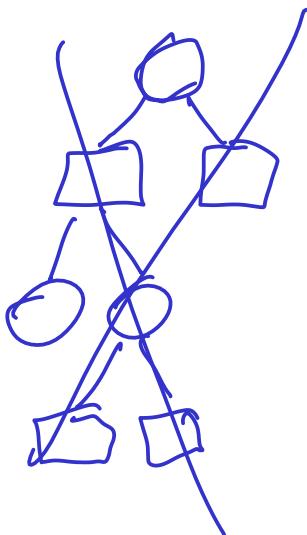
$$V(s) = \max_{a \in \mathcal{A}} (R(s, a) + \mathbb{E}[V(s')])$$

Minimax Tree

Minimax Trees

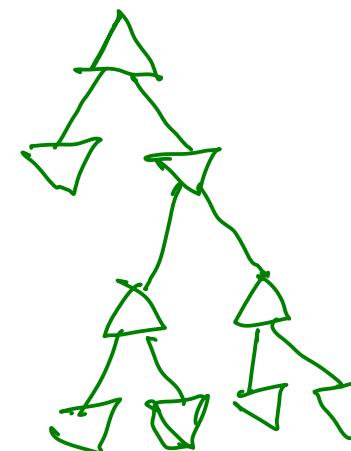
MDP Expectimax Tree

$$V(s) = \max_{a \in \mathcal{A}} (R(s, a) + \mathbb{E}[V(s')])$$



Minimax Tree

$$V_1(s) = \max_{a \in \mathcal{A}_1} (R(s, a) + \min_{a' \in \mathcal{A}_2} (R(s', a') + V(s'')))$$



Tree Backup Example

Tree Backup Example

```
function MINIMAX-DECISION(state) returns an action
    return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(s, a))$ 
```

```
function MAX-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
     $v \leftarrow -\infty$ 
    for each a in ACTIONS(state) do
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$ 
    return v
```

```
function MIN-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
     $v \leftarrow \infty$ 
    for each a in ACTIONS(state) do
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$ 
    return v
```

Tree Backup Example

```
function MINIMAX-DECISION(state) returns an action
    return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(s, a))$ 



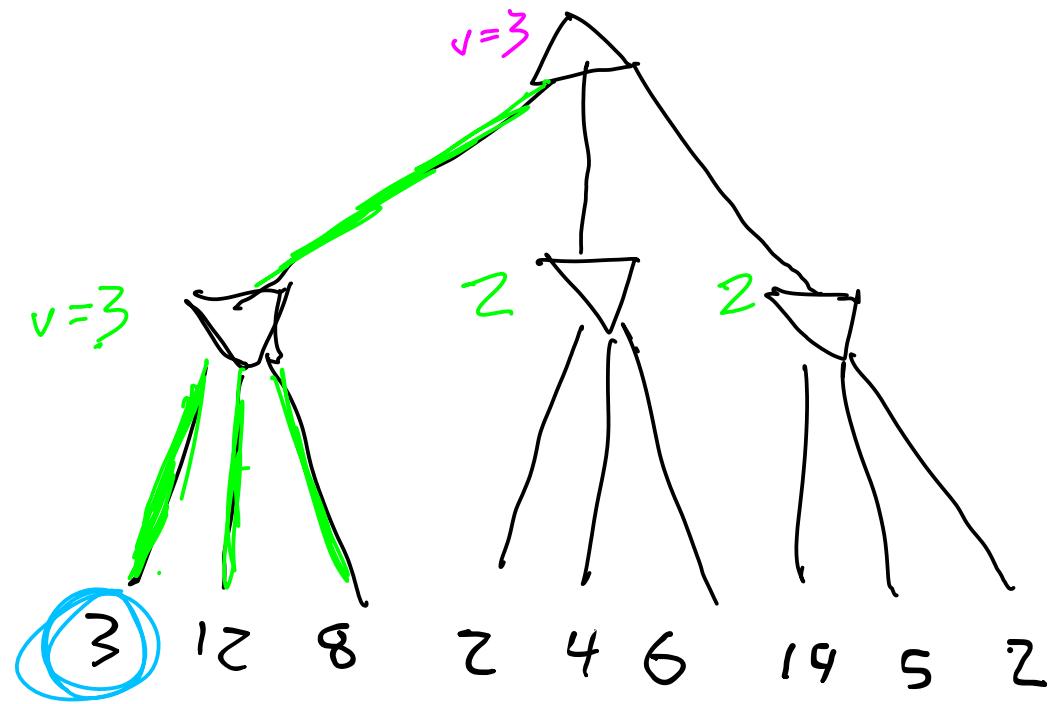
---

  
function MAX-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
    v  $\leftarrow -\infty$ 
    for each a in ACTIONS(state) do
        v  $\leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$ 
    return v



---

  
function MIN-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
    v  $\leftarrow \infty$ 
    for each a in ACTIONS(state) do
        v  $\leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$ 
    return v
```



Tree Backup Example

```
function MINIMAX-DECISION(state) returns an action
    return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(s, a))$ 
```

```
function MAX-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
     $v \leftarrow -\infty$ 
    for each a in ACTIONS(state) do
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$ 
    return v
```

```
function MIN-VALUE(state) returns a utility value
    if TERMINAL-TEST(state) then return UTILITY(state)
     $v \leftarrow \infty$ 
    for each a in ACTIONS(state) do
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$ 
    return v
```

3 12 8 2 4 6 19 5 2

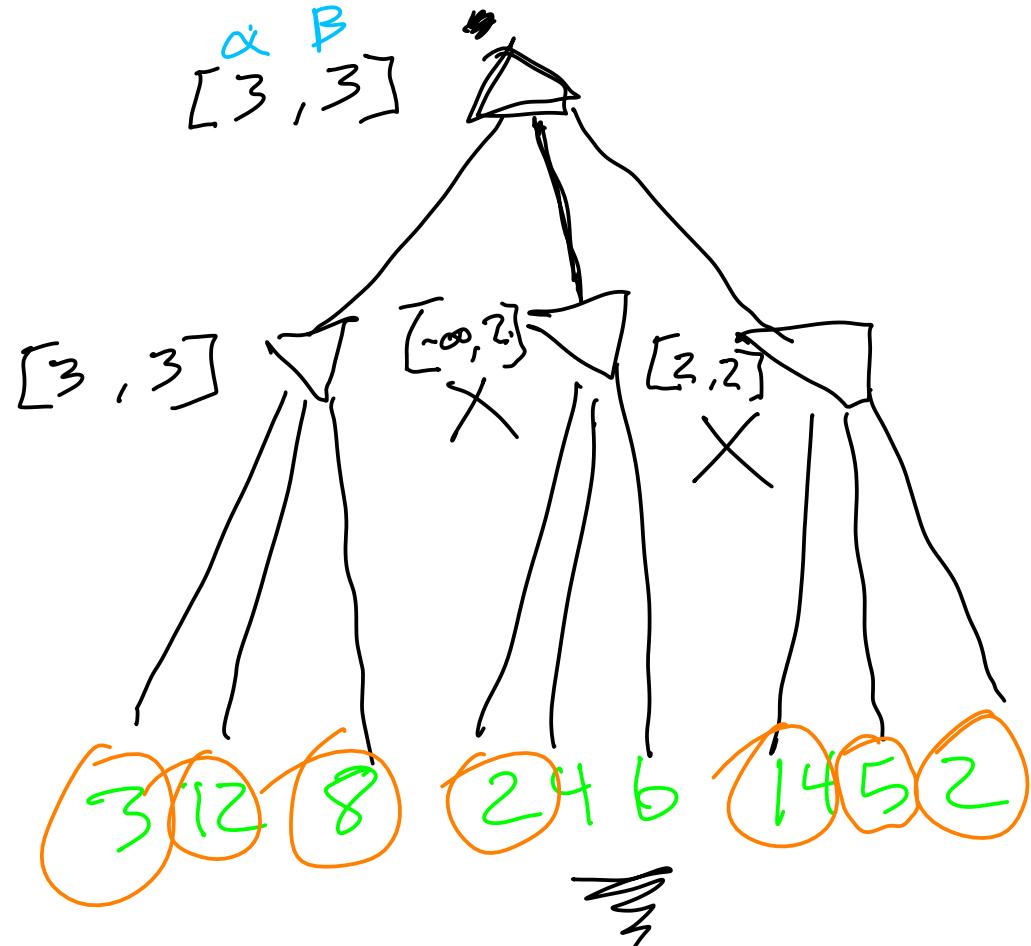
Why is this harder than an MDP? (think back to sparse sampling)

Alpha-Beta Pruning

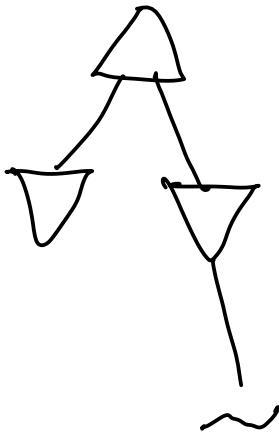
```
function ALPHA-BETA-SEARCH(state) returns an action
  v  $\leftarrow$  MAX-VALUE(state,  $-\infty$ ,  $+\infty$ )
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v  $\leftarrow -\infty$ 
  for each a in ACTIONS(state) do
    v  $\leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s,a), \alpha, \beta))$ 
    if v  $\geq \beta$  then return v
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v  $\leftarrow +\infty$ 
  for each a in ACTIONS(state) do
    v  $\leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s,a), \alpha, \beta))$ 
    if v  $\leq \alpha$  then return v
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return v
```



Evaluation Functions



Evaluation Functions

Examples of evaluation functions:

Evaluation Functions

Examples of evaluation functions:

- Rollout using default policy

Evaluation Functions

Examples of evaluation functions:

- Rollout using default policy
- Neural Network

Evaluation Functions

Examples of evaluation functions:

- Rollout using default policy
- Neural Network
- Weighted combination of features (e.g. number of Pawns, Knights, etc.)

Deep Blue



Deep Blue

- Defeated world champion Gary Kasparov in may 1997



Deep Blue

- Defeated world champion Gary Kasparov in may 1997
- 30 IBM RS/600 processors for alpha-beta search



Deep Blue

- Defeated world champion Gary Kasparov in may 1997
- 30 IBM RS/600 processors for alpha-beta search
- Routinely reached depth 14, sometimes up to 40



Deep Blue

- Defeated world champion Gary Kasparov in may 1997
- 30 IBM RS/600 processors for alpha-beta search
- Routinely reached depth 14, sometimes up to 40
- Evaluation function with 8000 features



Deep Blue

- Defeated world champion Gary Kasparov in may 1997
- 30 IBM RS/600 processors for alpha-beta search
- Routinely reached depth 14, sometimes up to 40
- Evaluation function with 8000 features
- Opening book of 4000 positions

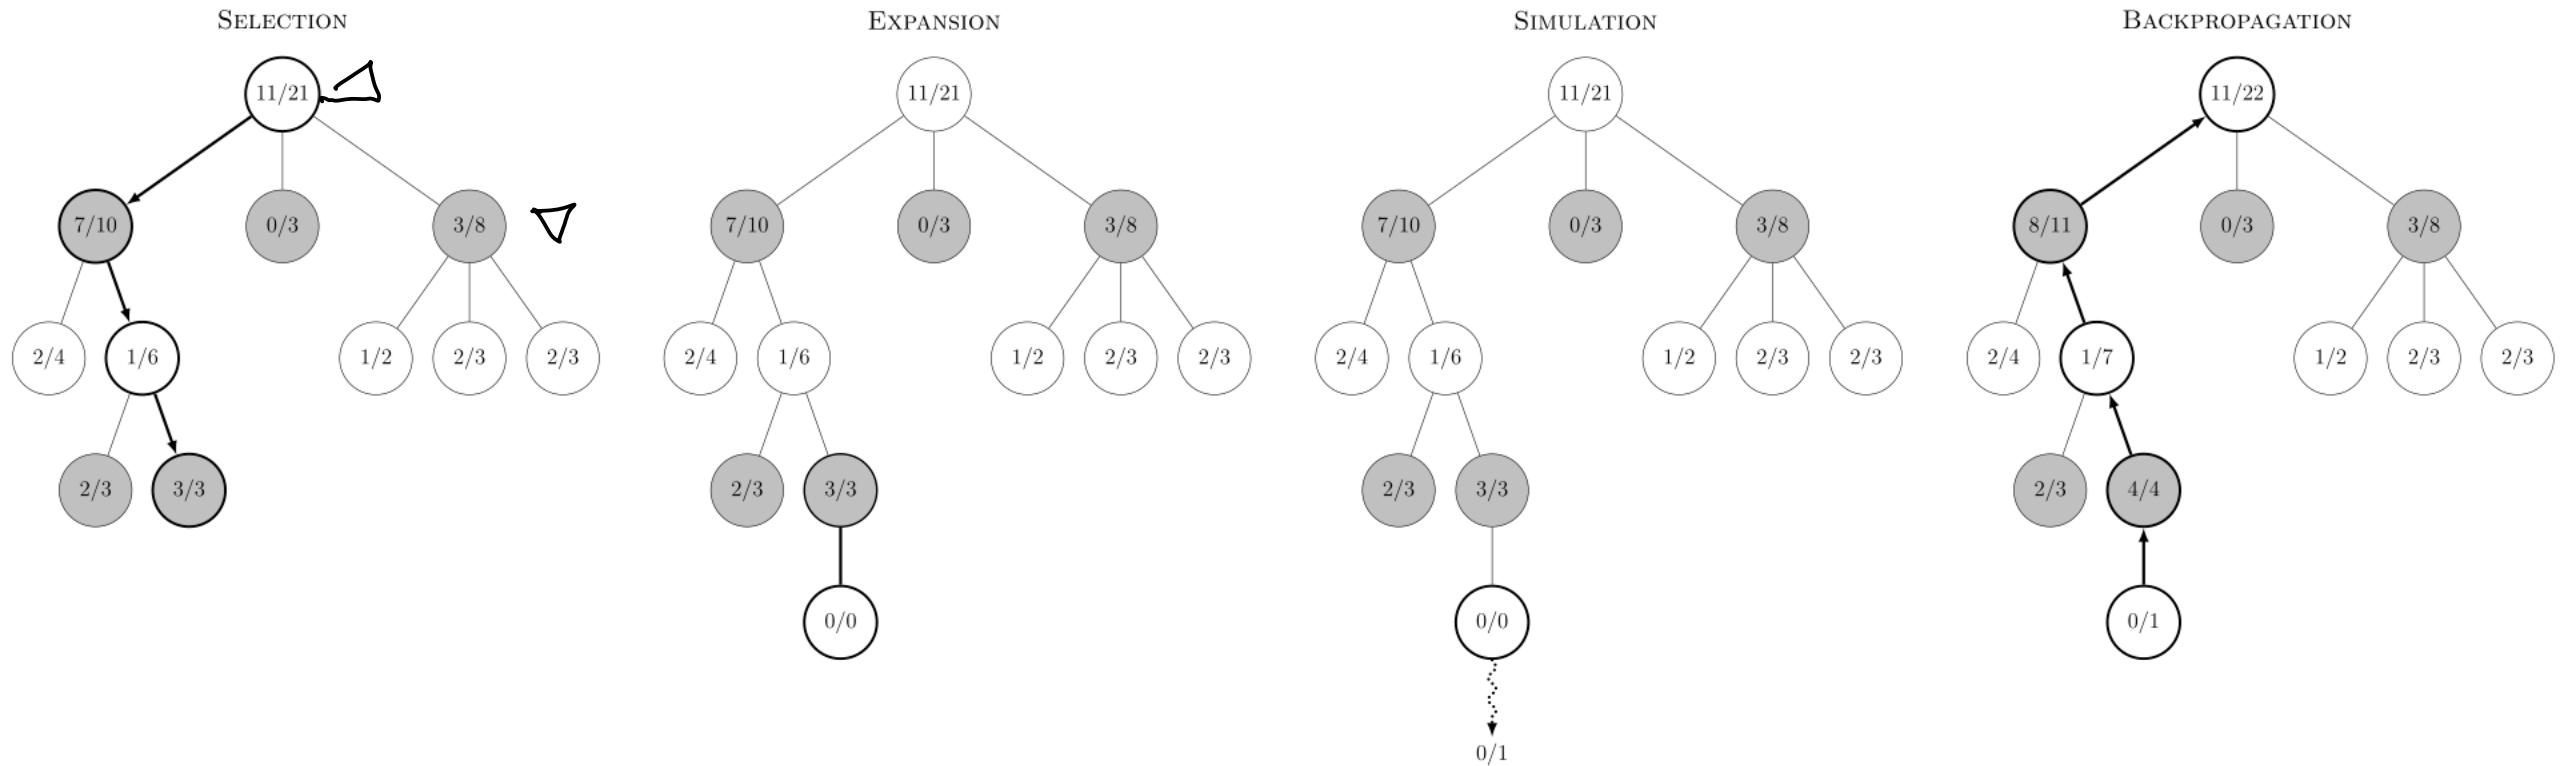


Deep Blue

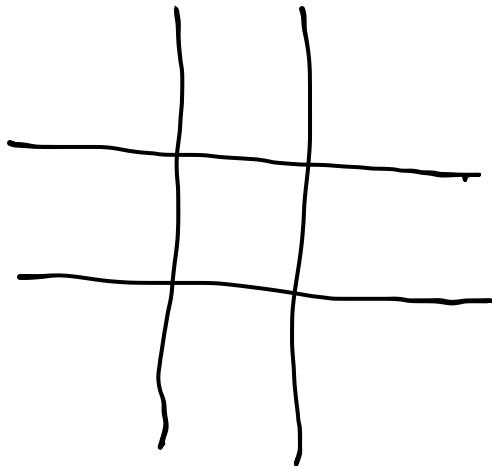
- Defeated world champion Gary Kasparov in may 1997
- 30 IBM RS/600 processors for alpha-beta search
- Routinely reached depth 14, sometimes up to 40
- Evaluation function with 8000 features
- Opening book of 4000 positions
- Database of 700,000 grandmaster games



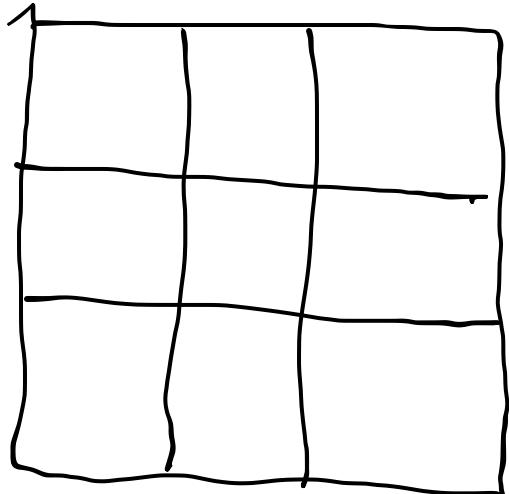
MCTS for Games



Break Exercise



Break Exercise



Break Exercise

1	2	3
4	5	6
7	8	9

Break Exercise

1	2	3
4	5	6
7	8	9

- Start at 5

Break Exercise

1	2	3
4	5	6
7	8	9

- Start at 5
- Max can move up/down

Break Exercise

1	2	3
4	5	6
7	8	9

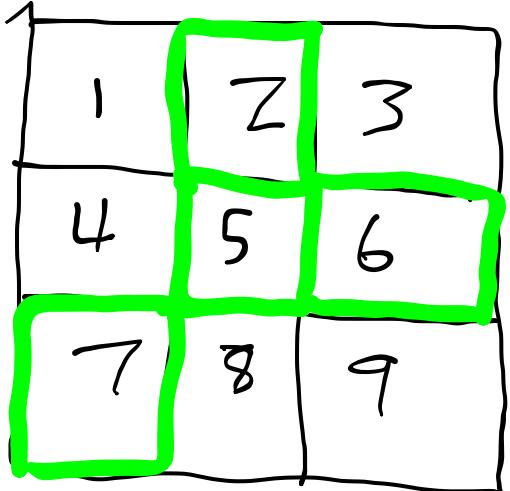
- Start at 5
- Max can move up/down
- Min moves left/right

Break Exercise

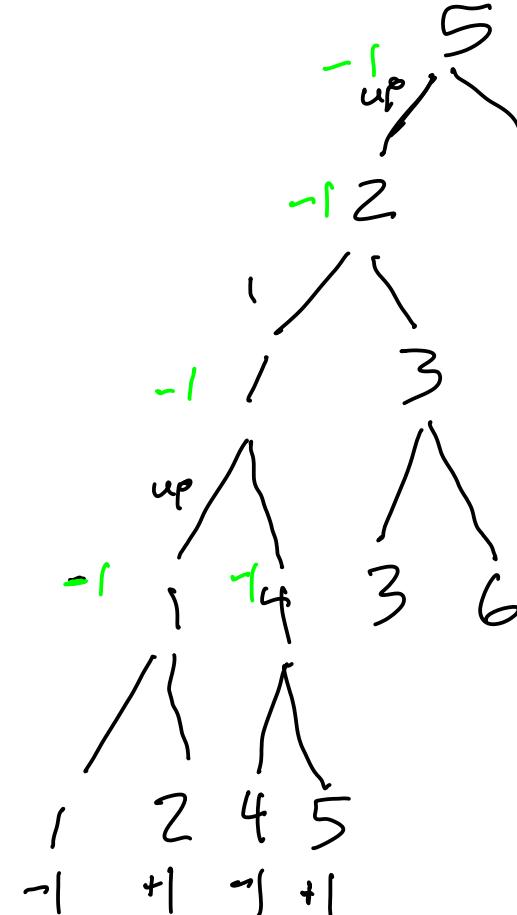
1	2	3
4	5	6
7	8	9

- Start at 5
- Max can move up/down
- Min moves left/right
- Each player gets 2 turns

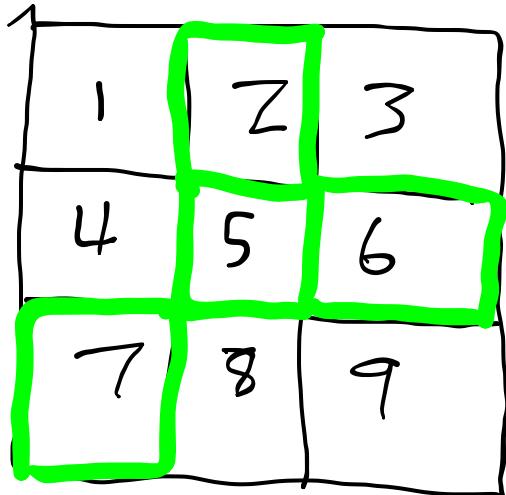
Break Exercise



- Start at 5
- Max can move up/down
- Min moves left/right
- Each player gets 2 turns
- On green squares, Max wins

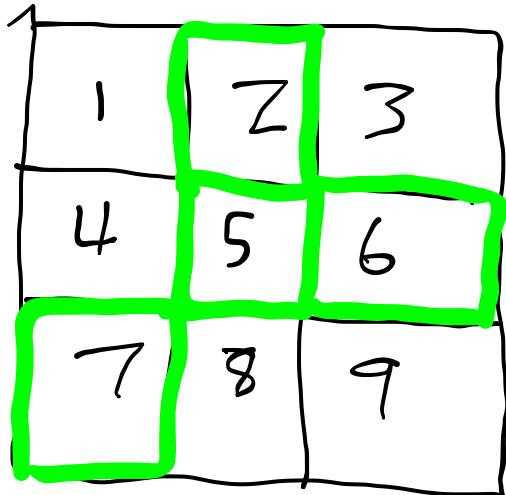


Break Exercise



- Start at 5
- Max can move up/down
- Min moves left/right
- Each player gets 2 turns
- On green squares, Max wins
- Who has advantage and what is first move?

Break Exercise



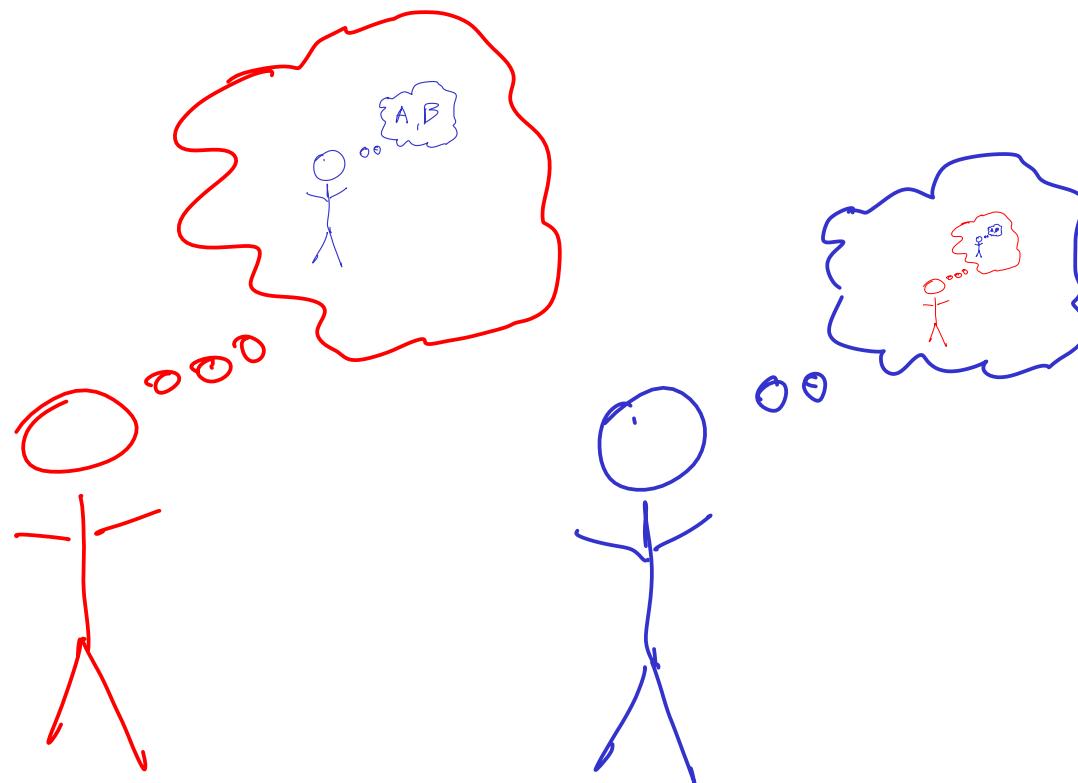
- Start at 5
- Max can move up/down
- Min moves left/right
- Each player gets 2 turns
- On green squares, Max wins
- Who has advantage and what is first move?

Bonus: what if 5 was not green?

Incomplete Information



Incomplete Information



Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

- Finite set of n players, plus the "chance" player

Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

- Finite set of n players, plus the "chance" player
- $P(h)$ (player at each history)

Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

- Finite set of n players, plus the "chance" player
- $P(h)$ (player at each history)
- $A(h)$ (set of actions at each history)

Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

- Finite set of n players, plus the "chance" player
- $P(h)$ (player at each history)
- $A(h)$ (set of actions at each history)
- $I(h)$ (information set that each history maps to)

Incomplete Information

Extensive-form game definition (h is a sequence of actions called a "history"):

- Finite set of n players, plus the "chance" player
- $P(h)$ (player at each history)
- $A(h)$ (set of actions at each history)
- $I(h)$ (information set that each history maps to)
- $U(h)$ (payoff for each leaf node in the game tree)

King-Ace Poker Example

King-Ace Poker Example

- 4 Cards: 2 Aces, 2 Kings

King-Ace Poker Example

- 4 Cards: 2 Aces, 2 Kings
- Each player is dealt a card

King-Ace Poker Example

- 4 Cards: 2 Aces, 2 Kings
- Each player is dealt a card
- P1 can either *raise* (r) the payoff
to 2 points or *check* (k) the
payoff at 1 point

King-Ace Poker Example

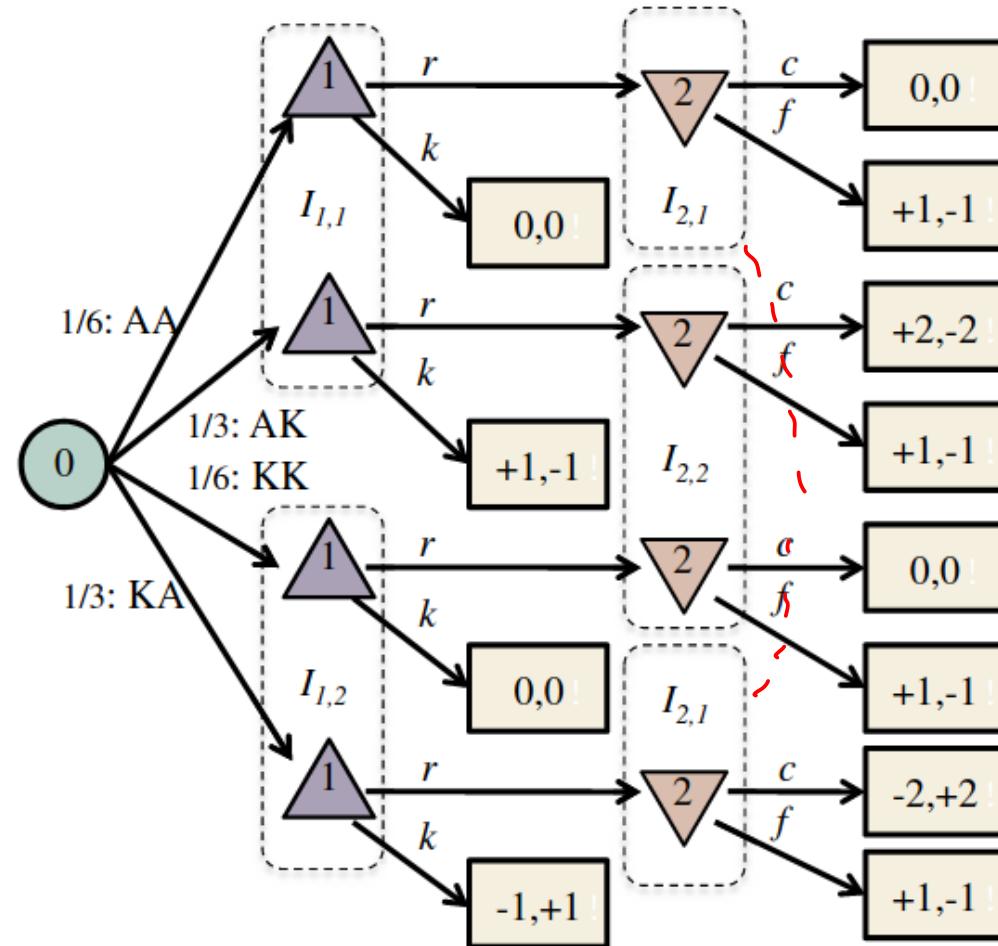
- 4 Cards: 2 Aces, 2 Kings
- Each player is dealt a card
- P1 can either *raise* (r) the payoff to 2 points or *check* (k) the payoff at 1 point
- If P1 raises, P2 can either *call* (c) Player 1's bet, or *fold* (f) the payoff back to 1 point

King-Ace Poker Example

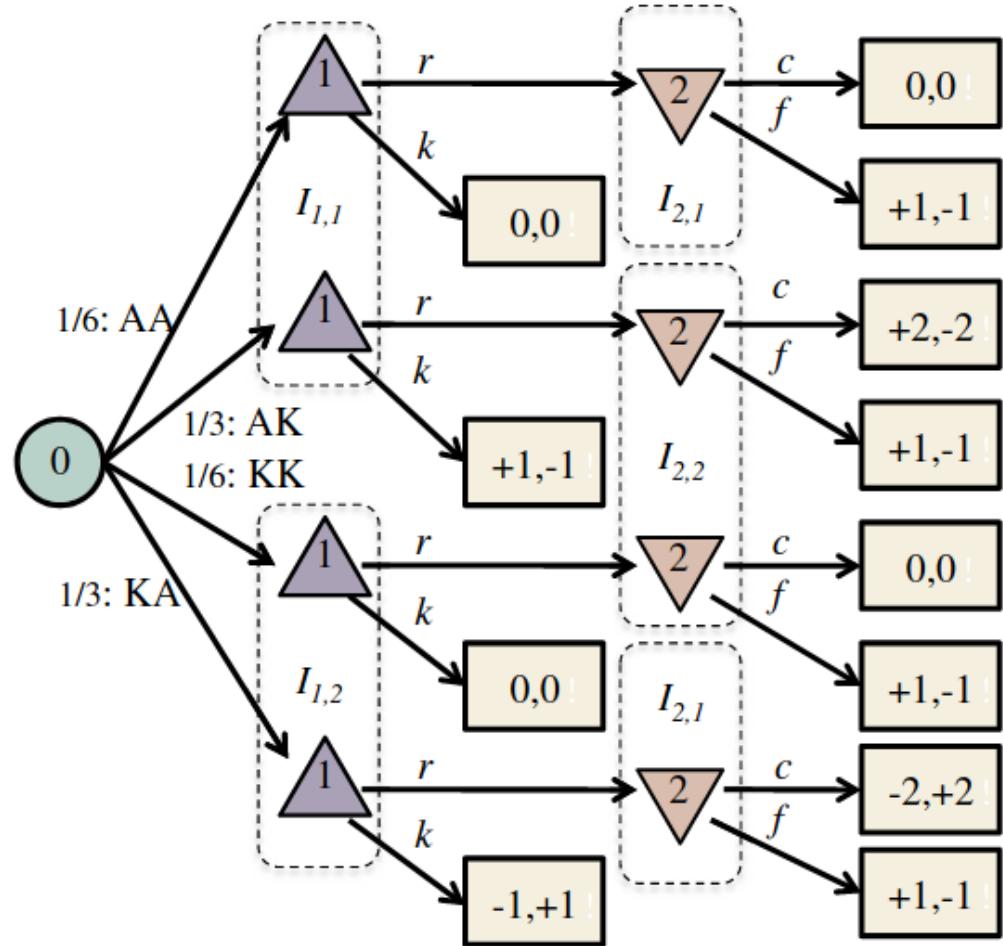
- 4 Cards: 2 Aces, 2 Kings
- Each player is dealt a card
- P1 can either *raise* (r) the payoff to 2 points or *check* (k) the payoff at 1 point
- If P1 raises, P2 can either *call* (c) Player 1's bet, or *fold* (f) the payoff back to 1 point
- The highest card wins

King-Ace Poker Example

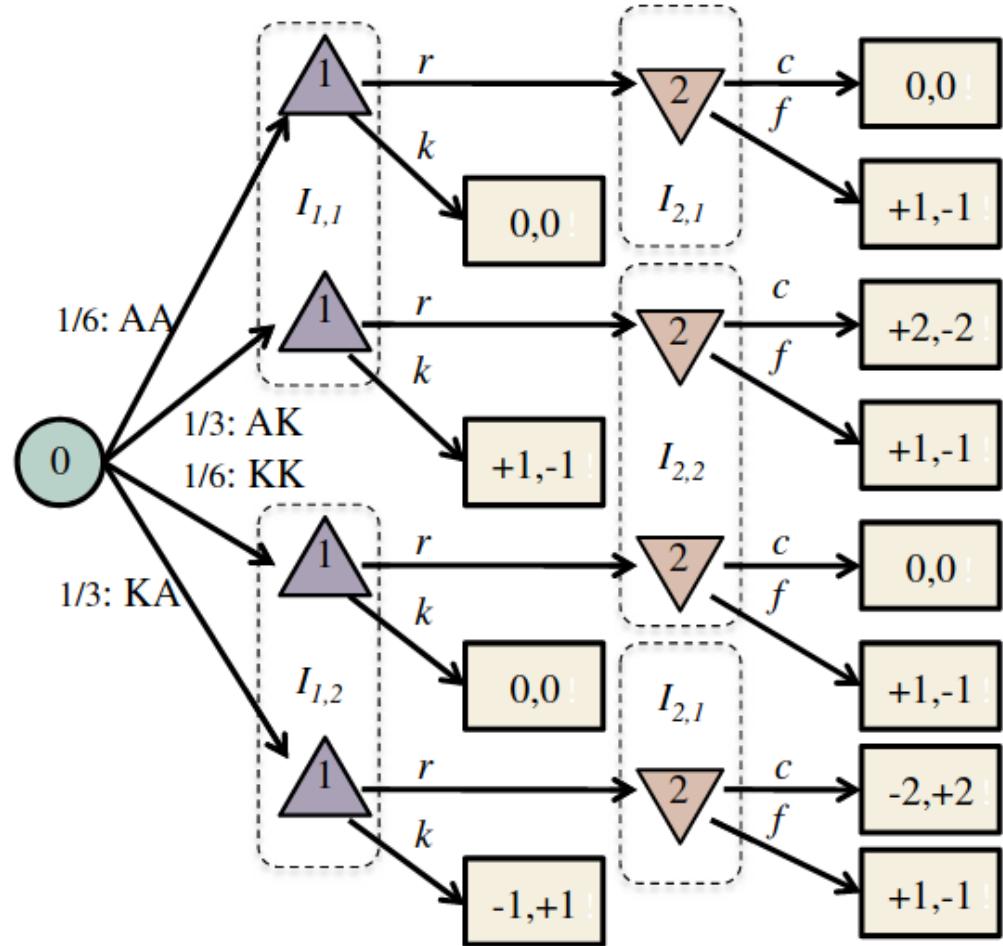
- 4 Cards: 2 Aces, 2 Kings
- Each player is dealt a card
- P1 can either *raise* (r) the payoff to 2 points or *check* (k) the payoff at 1 point
- If P1 raises, P2 can either *call* (c) Player 1's bet, or *fold* (f) the payoff back to 1 point
- The highest card wins



Extensive to Matrix Form

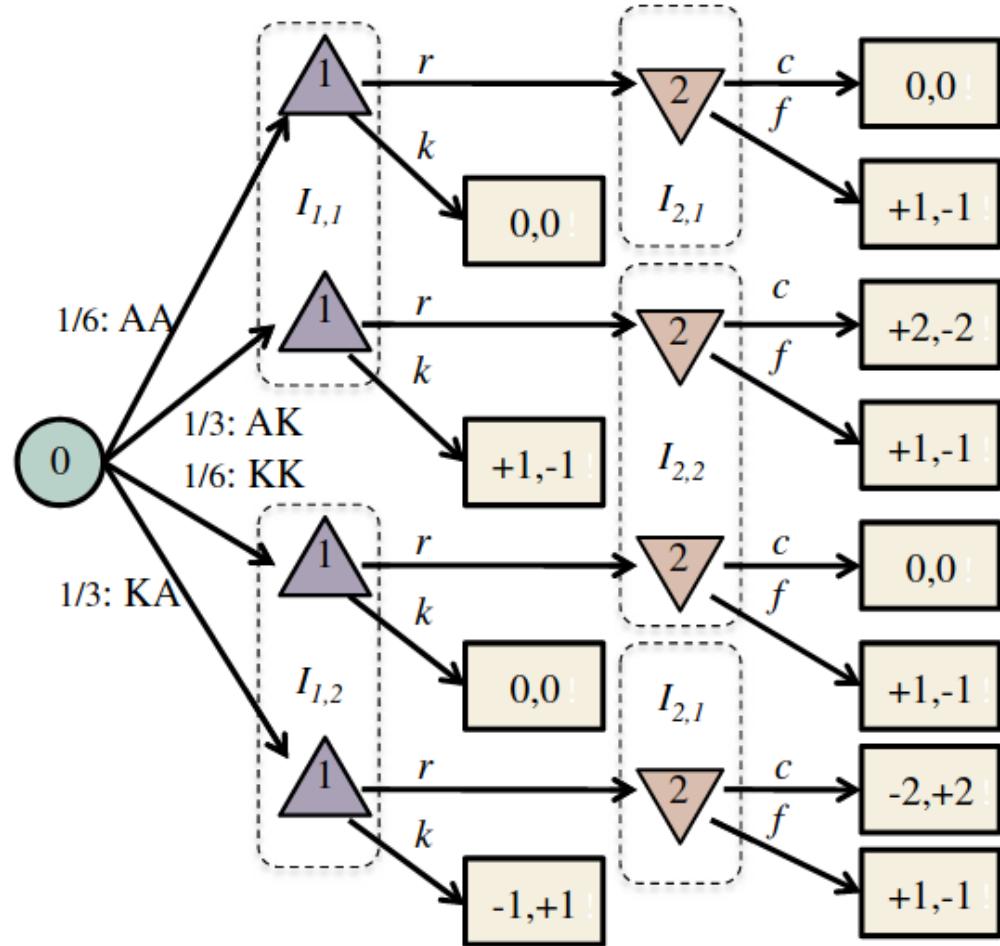


Extensive to Matrix Form



	2:cc	2:cf	2:ff	2:fc
1:rr	0	-1/6	1	7/6
1:kr	-1/3	-1/6	5/6	2/3
1:rk	1/3	0	1/6	1/2
1:kk	0	0	0	0

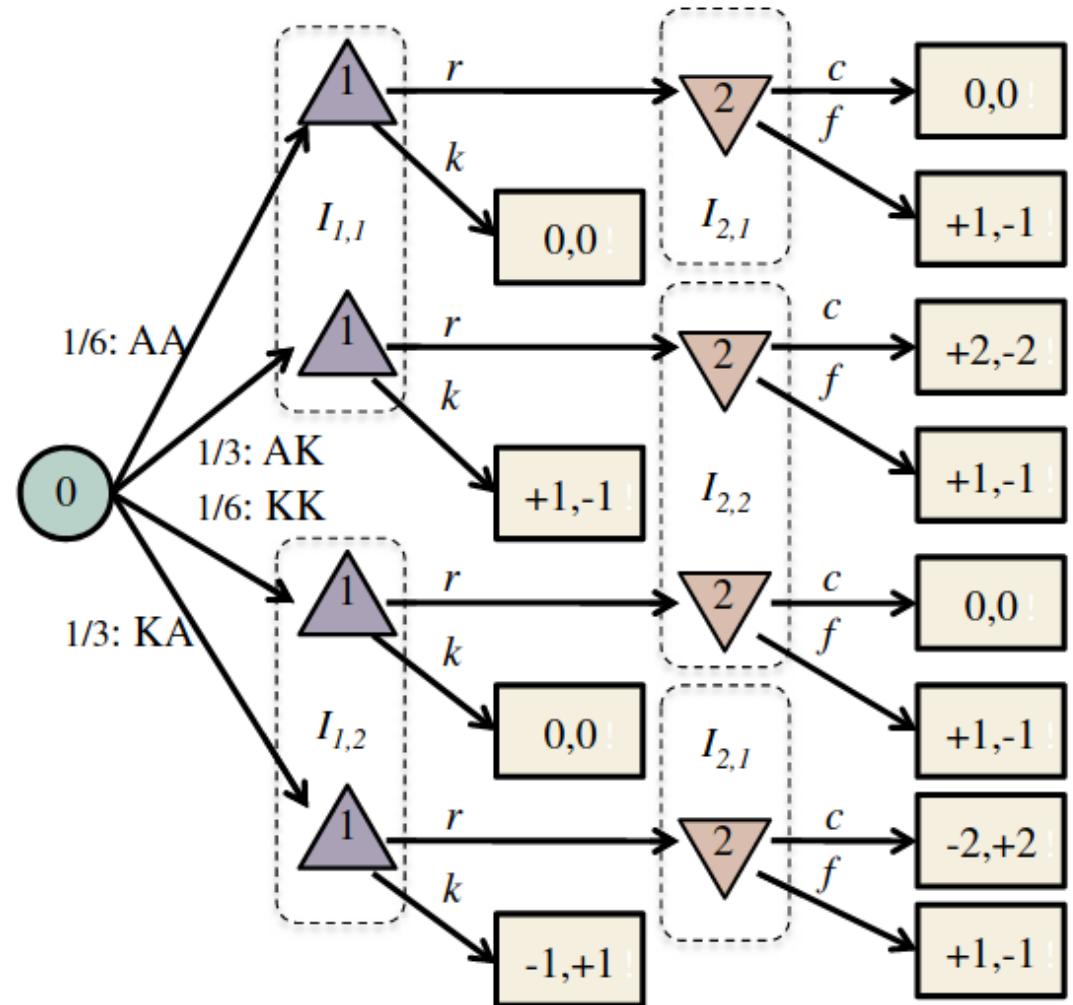
Extensive to Matrix Form



	2:cc	2:cf	2:ff	2:fc
1:rr	0	-1/6	1	7/6
1:kr	-1/3	-1/6	5/6	2/3
1:rk	1/3	0	1/6	1/2
1:kk	0	0	0	0

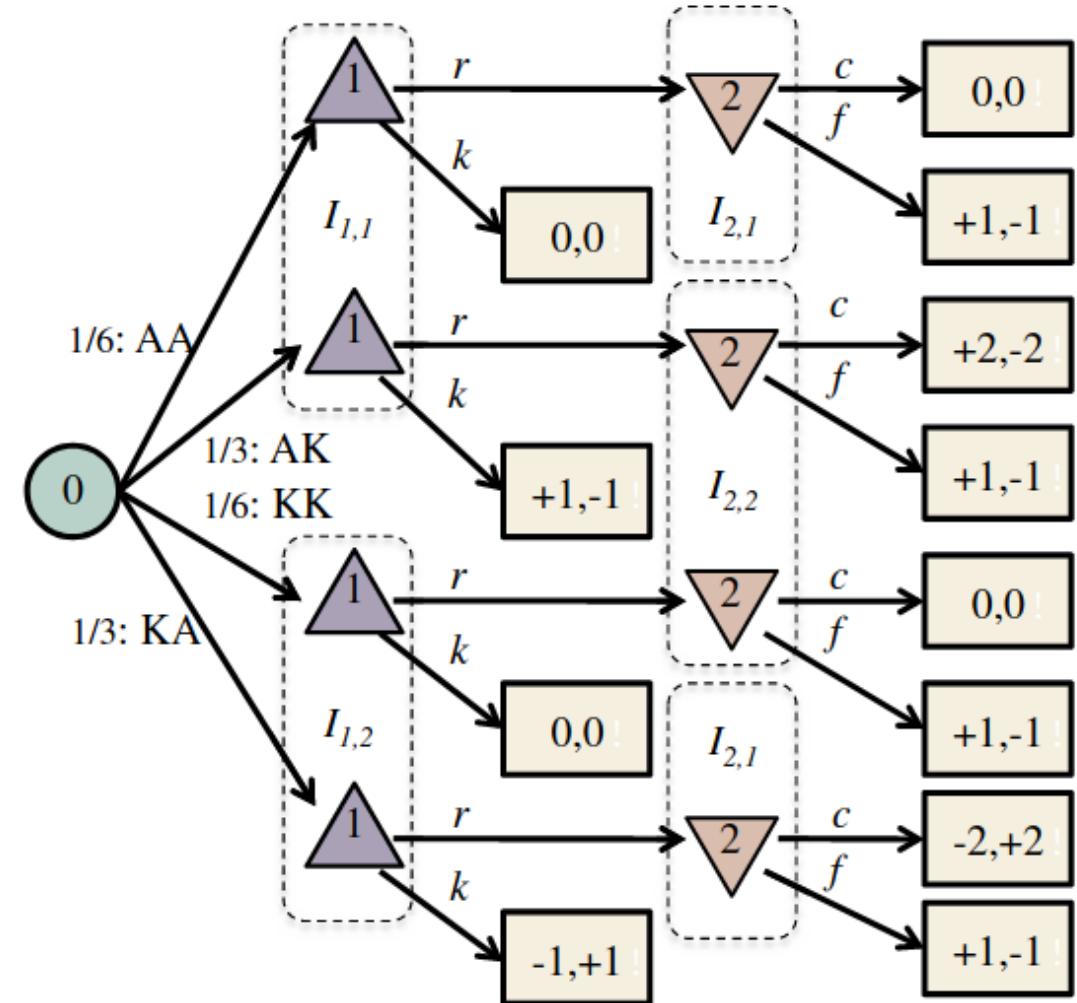
Exponential in number of info states!

Counterfactual Regret Minimization



Counterfactual Regret Minimization

- 1: Initialize cumulative regret tables: $\forall I, r_I[a] \leftarrow 0$.
- 2: Initialize cumulative strategy tables: $\forall I, s_I[a] \leftarrow 0$.
- 3: Initialize initial profile: $\sigma^1(I, a) \leftarrow 1/|A(I)|$



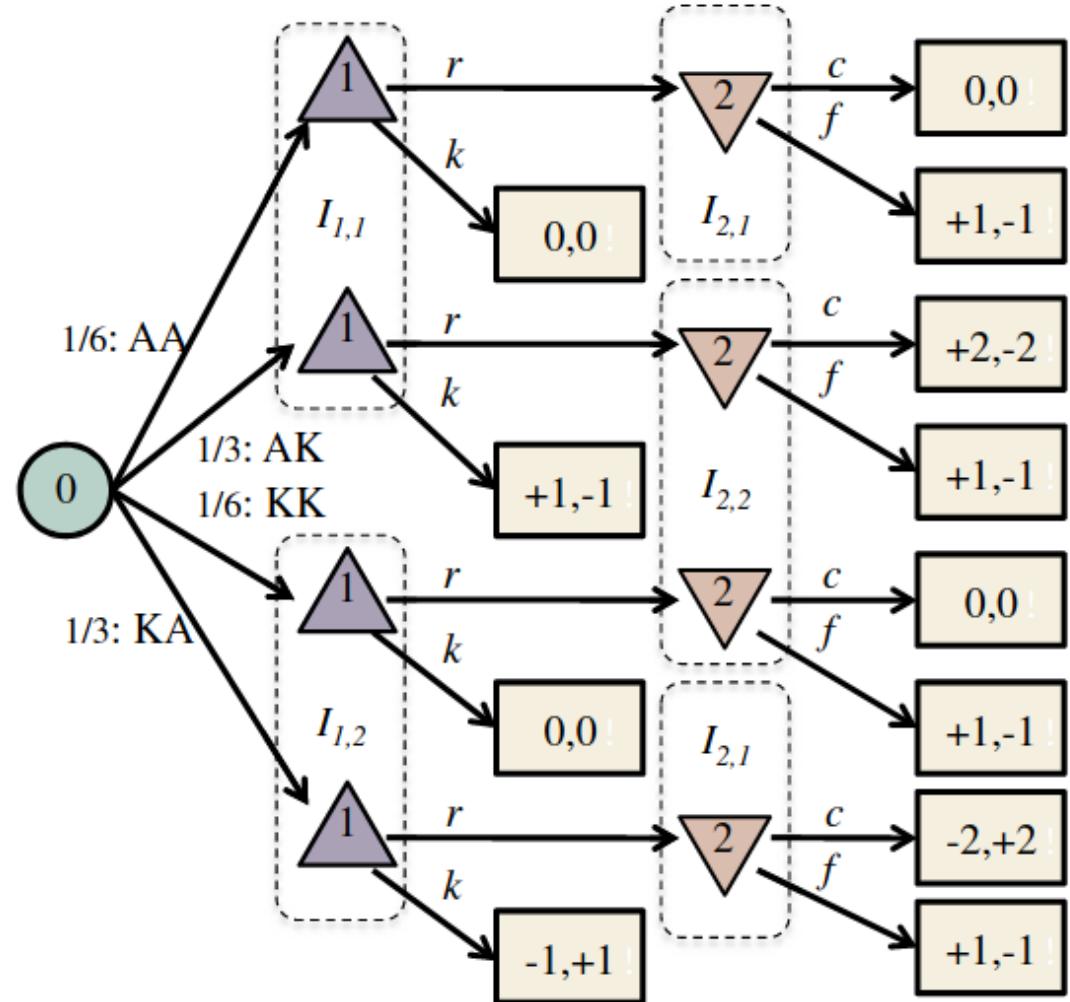
Counterfactual Regret Minimization

- 1: Initialize cumulative regret tables: $\forall I, r_I[a] \leftarrow 0$.
- 2: Initialize cumulative strategy tables: $\forall I, s_I[a] \leftarrow 0$.
- 3: Initialize initial profile: $\sigma^1(I, a) \leftarrow 1/|A(I)|$

```

32: function Solve():
33:   for  $t = \{1, 2, 3, \dots, T\}$  do
34:     for  $i \in \{1, 2\}$  do
35:       CFR( $\emptyset, i, t, 1, 1$ )
36:     end for
37:   end for

```



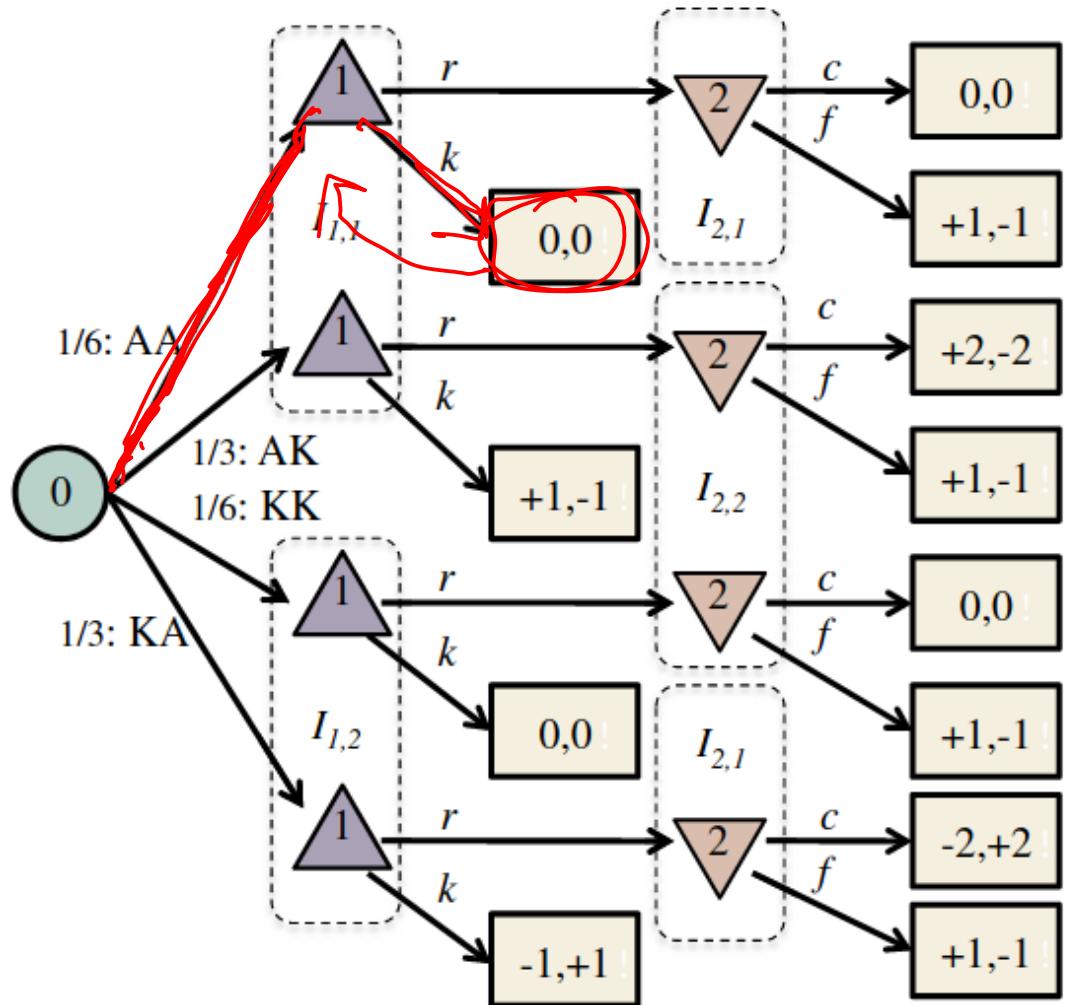
Counterfactual Regret Minimization

```

1: Initialize cumulative regret tables:  $\forall I, r_I[a] \leftarrow 0$ .
2: Initialize cumulative strategy tables:  $\forall I, s_I[a] \leftarrow 0$ .
3: Initialize initial profile:  $\sigma^1(I, a) \leftarrow 1/|A(I)|$ 

5: function CFR( $h, i, t, \pi_1, \pi_2$ ):
6:   if  $h$  is terminal then
7:     return  $u_i(h)$ 
8:   else if  $h$  is a chance node then
9:     Sample a single outcome  $a \sim \sigma_c(h, a)$ 
10:    return CFR( $ha, i, t, \pi_1, \pi_2$ )
11:   end if
12:   Let  $I$  be the information set containing  $h$ .
13:    $v_\sigma \leftarrow 0$ 
14:    $v_{\sigma_{I \rightarrow a}}[a] \leftarrow 0$  for all  $a \in A(I)$ 
15:   for  $a \in A(I)$  do
16:     if  $P(h) = 1$  then
17:        $v_{\sigma_{I \rightarrow a}}[a] \leftarrow \text{CFR}(ha, i, t, \sigma^t(I, a) \cdot \pi_1, \pi_2)$ 
18:     else if  $P(h) = 2$  then
19:        $v_{\sigma_{I \rightarrow a}}[a] \leftarrow \text{CFR}(ha, i, t, \pi_1, \sigma^t(I, a) \cdot \pi_2)$ 
20:     end if
21:      $v_\sigma \leftarrow v_\sigma + \sigma^t(I, a) \cdot v_{\sigma_{I \rightarrow a}}[a]$ 
22:   end for
23:   if  $P(h) = i$  then
24:     for  $a \in A(I)$  do
25:        $r_I[a] \leftarrow r_I[a] + \pi_{-i} \cdot (v_{\sigma_{I \rightarrow a}}[a] - v_\sigma)$ 
26:        $s_I[a] \leftarrow s_I[a] + \pi_i \cdot \sigma^t(I, a)$ 
27:     end for
28:      $\sigma^{t+1}(I) \leftarrow \text{regret-matching values}$ 
29:   end if
30:   return  $v_\sigma$ 

```



Kuhn Poker Example

3 Cards: 1, 2, and 3

Sequential Actions			Payoff
Player 1	Player 2	Player 1	
pass	pass		+1 to player with higher card
pass	bet	pass	+1 to player 2
pass	bet	bet	+2 to player with higher card
bet	pass		+1 to player 1
bet	bet		+2 to player with higher card

Kuhn Poker Example

3 Cards: 1, 2, and 3

Sequential Actions			Payoff
Player 1	Player 2	Player 1	
pass	pass		+1 to player with higher card
pass	bet	pass	+1 to player 2
pass	bet	bet	+2 to player with higher card
bet	pass		+1 to player 1
bet	bet		+2 to player with higher card

Solution (found by hand in 1950s):

- Player 1 may pass on a 3 with arbitrary probability γ , but then they must bet on a 1 w.p. $\gamma/3$ and bet on a 2 in the second round w.p. $\gamma/3 + 1/3$
- Player 2 must bet 1/3 of the time when holding a 1 after a pass and bet 1/3 of the time when holding a 2 with a bet

Matrix to Extensive Form

Matrix to Extensive Form

	R	P	S
R	$0, 0$	$-1, 1$	$1, -1$
P	$1, -1$	$0, 0$	$-1, 1$
S	$-1, 1$	$1, -1$	$0, 0$

Takeaways

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria
- Regret matching is a tool for finding Nash equilibria in Matrix games

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria
- Regret matching is a tool for finding Nash equilibria in Matrix games
- Turn-taking zero-sum games of perfect information can be solved using minimax trees with Alpha-beta pruning or MCTS

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria
- Regret matching is a tool for finding Nash equilibria in Matrix games
- Turn-taking zero-sum games of perfect information can be solved using minimax trees with Alpha-beta pruning or MCTS
- Imperfect information can be represented as an extensive-form game with strategies conditioned on information sets.

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria
- Regret matching is a tool for finding Nash equilibria in Matrix games
- Turn-taking zero-sum games of perfect information can be solved using minimax trees with Alpha-beta pruning or MCTS
- Imperfect information can be represented as an extensive-form game with strategies conditioned on information sets.
- Strategies for imperfect information games often involve stochastic "bluffing"

Takeaways

- Games do not have a single optimal solution; Instead they have equilibria
- Regret matching is a tool for finding Nash equilibria in Matrix games
- Turn-taking zero-sum games of perfect information can be solved using minimax trees with Alpha-beta pruning or MCTS
- Imperfect information can be represented as an extensive-form game with strategies conditioned on information sets.
- Strategies for imperfect information games often involve stochastic "bluffing"
- Counterfactual regret minimization can be used to solve extensive-form games